

CS 4320 / 7320


Advanced Database

Concepts

Maintenance

What is the SDLC?

Where does Maintenance fit?



Data Lakes and Data Warehouses

PSDS 3100
Database Analytics

Fifty Years of Database Technology in a Nutshell

Topics

Relational Databases

Entities

Relationships

OLAP Versus OLTP

Transactions

Analytics

Cluster Computing

Big Data

Spark

NOSQL

Graph Databases

Aggregation is Out

Connection is In

What is a database?

Wikipedia says:

“A database is an organized collection of data. The data are typically organized to model relevant aspects of reality in a way that supports processes requiring this information.”

Database Design

- What are the data entities of the system?
- What are the attributes of each entity?
- What are the constraints on the attributes of the entities?
- What can be used to uniquely identify entities in the system?
- How are the different entities related?

Entities

A person, place or thing about which we want to collect and store multiple instances of data.

Similar to an Object in Object Oriented Design

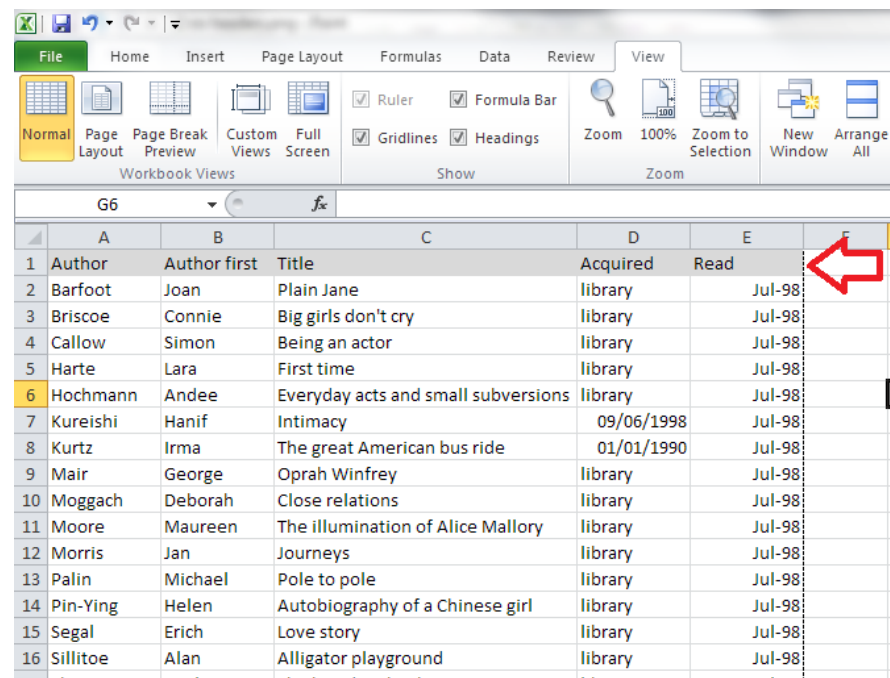
Think of Entities as nouns

These will be the tables in your database

Attributes

Data that describes the Entities.
These will be the columns of each table in your database.

Familiar concept?



The screenshot shows the Microsoft Excel 2003 interface. The 'View' tab is selected in the menu bar. The 'Workbook Views' group includes 'Normal' (selected), 'Page Layout', 'Page Break Preview', 'Custom Views', and 'Full Screen'. The 'Show' group includes 'Ruler', 'Gridlines', 'Formula Bar', and 'Headings'. The 'Zoom' group includes 'Zoom', '100%', 'Zoom to Selection', 'New Window', and 'Arrange All'. The active cell is G6. The table below is a library catalog with columns A through F. A red arrow points to the empty cell in column F, row 1.

	A	B	C	D	E	F
1	Author	Author first	Title	Acquired	Read	
2	Barfoot	Joan	Plain Jane	library	Jul-98	
3	Briscoe	Connie	Big girls don't cry	library	Jul-98	
4	Callow	Simon	Being an actor	library	Jul-98	
5	Harte	Lara	First time	library	Jul-98	
6	Hochmann	Andee	Everyday acts and small subversions	library	Jul-98	
7	Kureishi	Hanif	Intimacy	09/06/1998	Jul-98	
8	Kurtz	Irma	The great American bus ride	01/01/1990	Jul-98	
9	Mair	George	Oprah Winfrey	library	Jul-98	
10	Moggach	Deborah	Close relations	library	Jul-98	
11	Moore	Maureen	The illumination of Alice Mallory	library	Jul-98	
12	Morris	Jan	Journeys	library	Jul-98	
13	Palin	Michael	Pole to pole	library	Jul-98	
14	Pin-Ying	Helen	Autobiography of a Chinese girl	library	Jul-98	
15	Segal	Erich	Love story	library	Jul-98	
16	Sillitoe	Alan	Alligator playground	library	Jul-98	

Constraints

- Specific rules for the Attributes.
- Make sure that the data is consistent.
- In SQL:
 - NOT NULL
 - UNIQUE
 - CHECK
 - DEFAULT

How to access a record?

Primary Key: attribute or combination of attributes that uniquely identify each row in the table.

Relationships

- Entities have some relationship to other entities in the system
- Illustrates an association between two entities.
- Cardinality Constraints:
 - Zero or More
 - One or More
 - One and only One
 - Zero or One

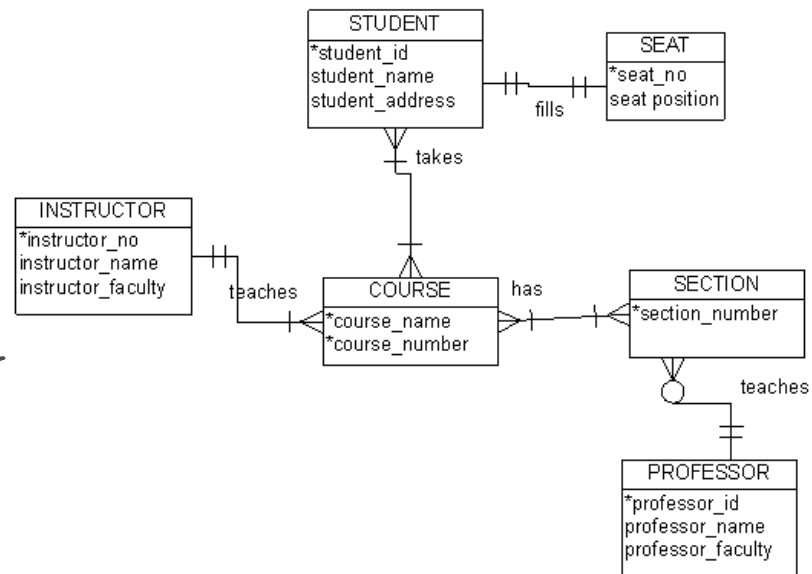
ERD: Entity Relationship Diagram

Entity-Relationship Diagram

- Data Model to describe a database in an abstract way
- Visually shows the relationships of data within the system
- We will use Crow's Foot Notation, because that is what real-world professionals use

ERD Example

Can I code
this exactly
as is??



```
CREATE TABLE Student (  
    student_id INT,  
    student_name VARCHAR(30),  
    student_address VARCHAR(40),  
    PRIMARY KEY (student_id)  
)
```

```
CREATE TABLE Course (  
    course_name VARCHAR(20),  
    course_number INT,  
    PRIMARY KEY (course_name, course_number)  
)
```

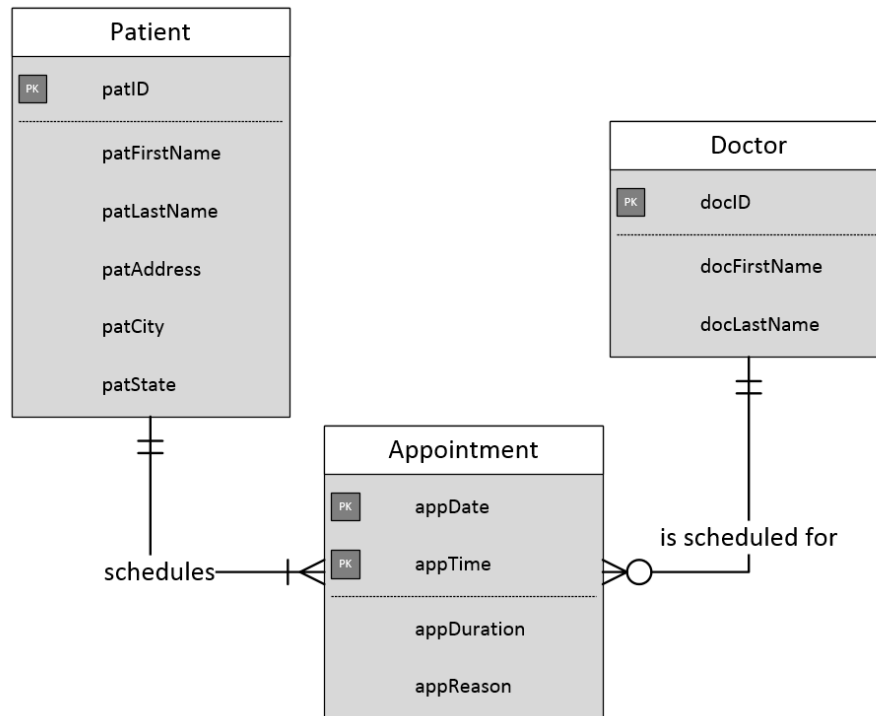


```
CREATE TABLE Enrollment (  
    student_id INT,  
    course_name VARCHAR(20),  
    course_number INT,  
    PRIMARY KEY  
        (student_id, course_name, course_number),  
    FOREIGN KEY (student_id)  
        REFERENCES Student(student_id),  
    FOREIGN KEY (course_name, course_number)  
        REFERENCES Course(course_name, course_number)  
)
```

Example

Think about a system that keeps track of doctor appointments scheduled by patients.

- What would the entities (tables) be?
- What would be the attributes (fields/columns) of each entity?
- What uniquely identifies the entities? (primary keys)
- How are the entities related?



Data Warehousing

Phase Two of Database Technology:

From Transactions to Analysis of Structured Data

Data Warehouses

Relational : Often Online Transaction Processing (OLTP)

- High volume inputs

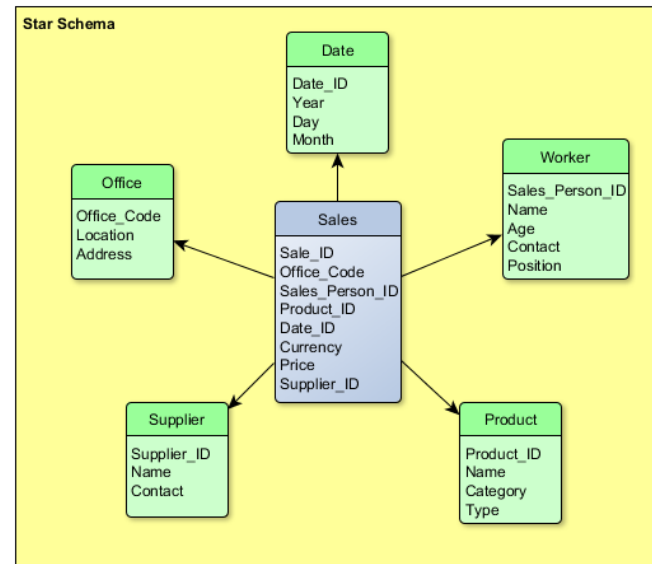
- Data organized to avoid redundancy

Data Warehouses: Often Online Analytical Processing (OLAP)

- Decision Support

- Designed for Analytical Speed

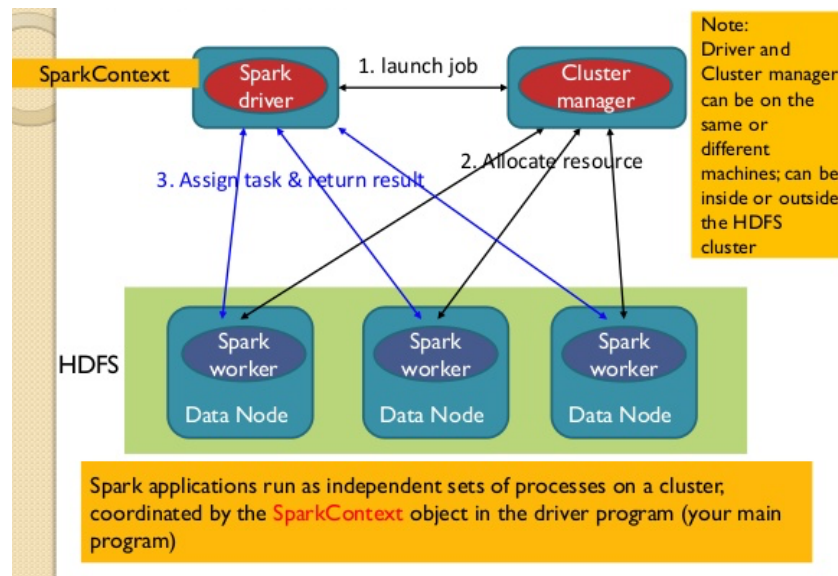
OLAP Database Design: Star Schema








Cluster Database Computing

Phase Three of Database Technology:
From Reorganizing and Moving Data to
Massive Parallel Processing

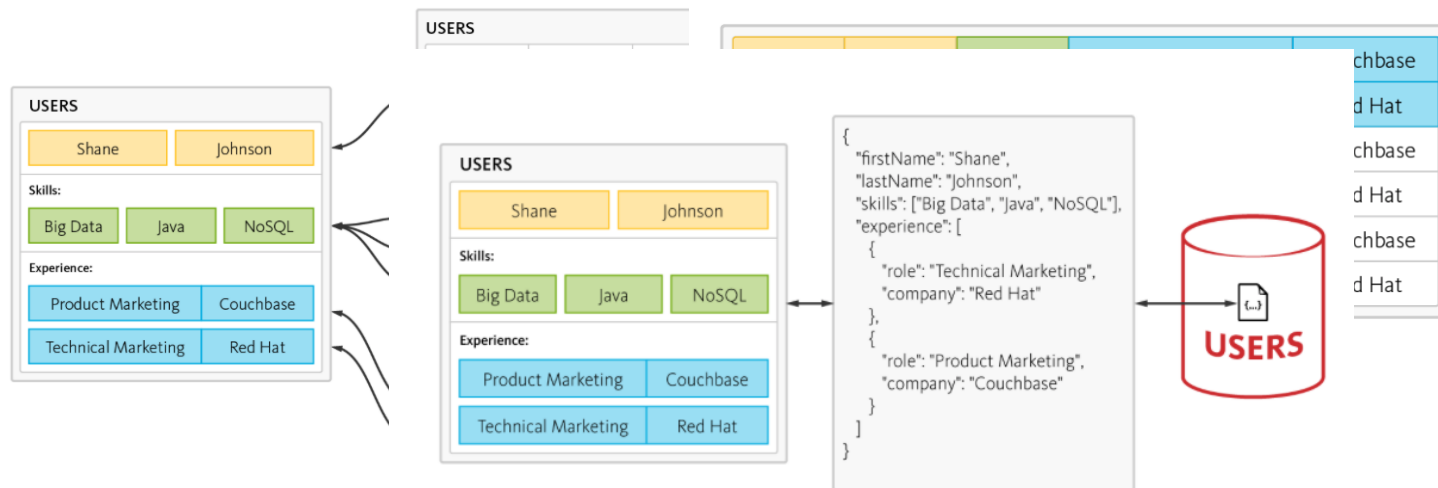
Big Data: Distributed Computing on Clusters



Big Data: Distributed Computing on Clusters

	 OPERATIONAL	ANALYTICAL 	
	 Couchbase	 Spark	 Hadoop
<i>Use cases</i>	<ul style="list-style-type: none">OperationalWeb / Mobile	<ul style="list-style-type: none">AnalyticsMachine Learning	<ul style="list-style-type: none">AnalyticsMachine Learning
<i>Processing mode</i>	<ul style="list-style-type: none">OnlineAd Hoc	<ul style="list-style-type: none">Ad HocBatchStreaming (+/-)	<ul style="list-style-type: none">BatchAd Hoc (+/-)
<i>Low latency =</i>	< 1 ms ops	Seconds	Minutes
<i>Performance</i>	Highly predictable	Variable	Variable
<i>Users are typically...</i>	Millions of customers	100's of analysts or data scientists	100's of analysts or data scientists
	Memory-centric	Memory-centric	Disk-centric
<i>Big data =</i>	10s of Terabytes	Petabytes	Petabytes

NOSQL



Graph Databases

Representing Everything as a Graph

Fifty Years of Database Technology

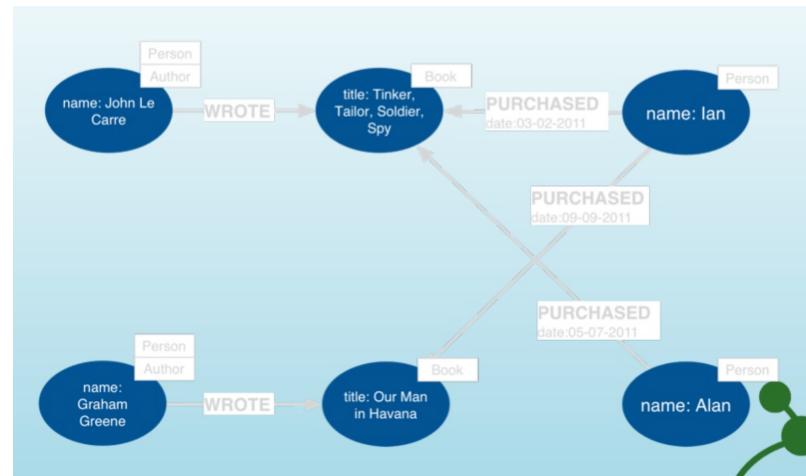
- OLTP / OLAP: Long History, Built on a relational model
- NOSQL / Spark: High Volume Inputs, Distributed Processing
- Graph Databases: Combining OLTP & OLAP Performance Characteristics in One Technology

Graph Databases

Aggregate v Connected Models: OLAP & OLTP are both “Aggregate” focused. Graph Databases are “Connection Focused”

Graph Building Blocks

- Nodes
 - Entities (tables)
- Relationships
 - Connect to Structure
- Properties
 - Attributes and metadata
- Labels
 - Group Nodes by Role

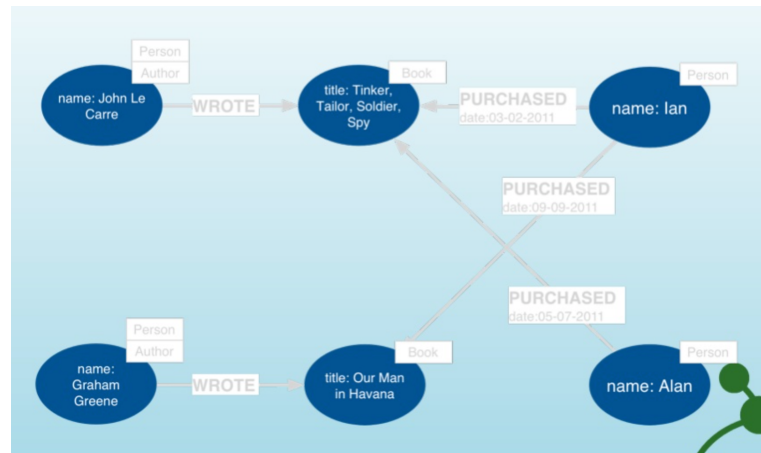


Graph Databases

Aggregate v Connected Models: OLAP & OLTP are both “Aggregate” focused. Graph Databases are “Connection Focused”

Graph Building Blocks

- Nodes
 - Entities (tables)
- Relationships
 - Connect to Structure
- Properties
 - Attributes and metadata
- Labels
 - Group Nodes by Role

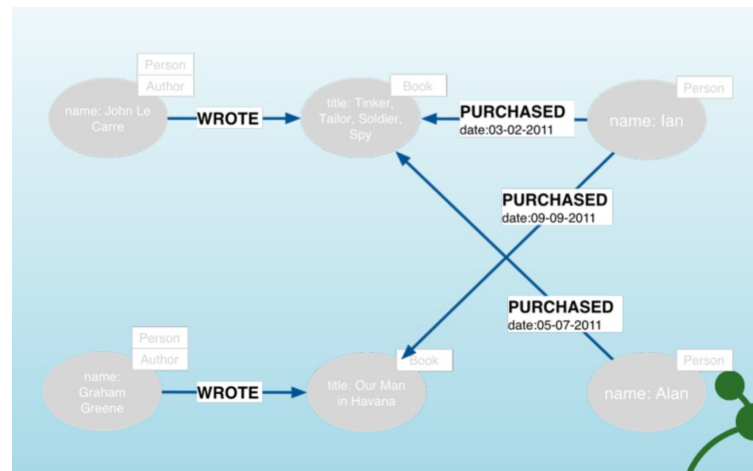


Graph Databases

Aggregate v Connected Models: OLAP & OLTP are both “Aggregate” focused. Graph Databases are “Connection Focused”

Graph Building Blocks

- Nodes
 - Entities (tables)
- Relationships
 - Connect to Structure
- Properties
 - Attributes and metadata
- Labels
 - Group Nodes by Role

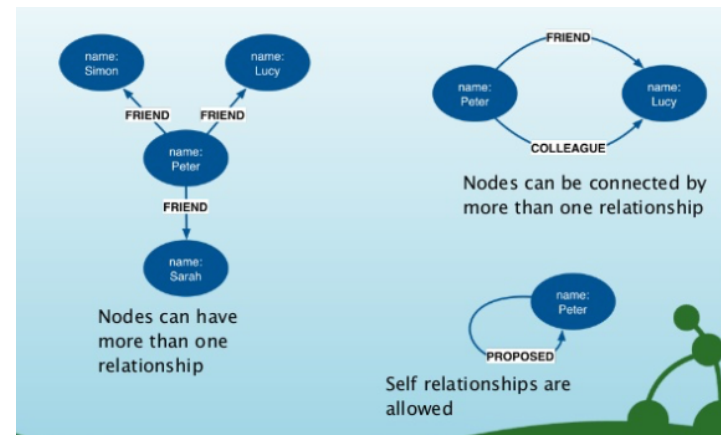


Graph Databases

Aggregate v Connected Models: OLAP & OLTP are both “Aggregate” focused. Graph Databases are “Connection Focused”

Graph Building Blocks

- Nodes
 - Entities (tables)
- Relationships
 - Connect to Structure
- Properties
 - Attributes and metadata
- Labels
 - Group Nodes by Role

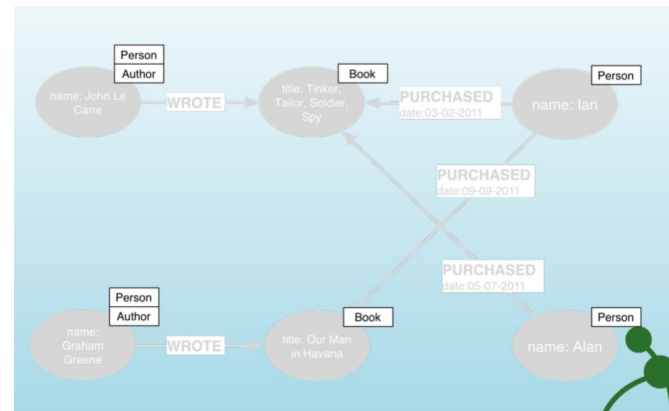


Graph Databases

Aggregate v Connected Models: OLAP & OLTP are both “Aggregate” focused. Graph Databases are “Connection Focused”

Graph Building Blocks

- Nodes
 - Entities (tables)
- Relationships
 - Connect to Structure
- Properties
 - Attributes and metadata
- Labels
 - Group Nodes by Role



Questions?